
```
import requests
import pandas as pd
from collections import Counter
import pandas as pd
import plotly.express as px
```

✓ Reading The Dataset

```
url = "https://api.github.com/repos/rails/rails/issues"

issues = []

params = {
    "state": "all",
    "per_page": 100,
}

for i in range(1, 6):
    params["page"] = i
    response = requests.get(url, params=params)
    issues.extend(response.json())

df = pd.DataFrame(issues)

df.head()
```

	url	repository_url	labels_url
0	https://api.github.com/repos/rails/rails/issue...	https://api.github.com/repos/rails/rails	https://api.github.com/repos/rails/rails/issue... https://api.github.co
1	https://api.github.com/repos/rails/rails/issue...	https://api.github.com/repos/rails/rails	https://api.github.com/repos/rails/rails/issue... https://api.github.co
2	https://api.github.com/repos/rails/rails/issue...	https://api.github.com/repos/rails/rails	https://api.github.com/repos/rails/rails/issue... https://api.github.co
3	https://api.github.com/repos/rails/rails/issue...	https://api.github.com/repos/rails/rails	https://api.github.com/repos/rails/rails/issue... https://api.github.co
4	https://api.github.com/repos/rails/rails/issue...	https://api.github.com/repos/rails/rails	https://api.github.com/repos/rails/rails/issue... https://api.github.co
5 rows × 30 columns			

✓ 1. How do the number of issues evolve over time?


```
df['created_at'] = pd.to_datetime(df['created_at'])

issue_counts_by_date = df.resample('D', on='created_at').size()

# Plot
fig = px.line(issue_counts_by_date, title='Number of Issues Over Time')
fig.update_xaxes(title_text='Date')
fig.update_yaxes(title_text='Number of Issues')
fig.show()
```



Number of Issues Over Time



2. Are there any periods in which we get more issues?

Yes, there are distinct periods where more issues are reported. These periods are represented by the peaks in the line chart. Specifically, from the provided image, it appears that:

- There is a significant peak in the number of issues reported around the first week of January 2024. This could indicate a surge in activity, possibly due to new year code sprints, releases, or other community activities. Another noticeable peak occurs in the third week of January 2024, suggesting another period with increased issue reporting.

- The final week of January 2024 also shows an increased number of issues, though not as pronounced as the first peak.
- These peaks could be due to various reasons such as new feature deployments, version updates, or discovery of bugs that coincide with these dates. It would be beneficial to cross-reference these dates with the Rails project's update logs, community forums, or other documentation to understand the context behind the increased number of issues.

✓ 3. Is there anyone who reports more issues than others?

```
df['reporter'] = df['user'].apply(lambda x: x['login'] if isinstance(x, dict) else None)
```

```
top_reporters = df['reporter'].value_counts().head(10)
```

```
print(top_reporters)
```

```
skipkayhil      33
seanpdoyle      26
p8              25
dhh             24
akhilkrishnan   21
casperisfine    14
sato11          14
byroot          11
Earlopain       10
dorianmarie     9
Name: reporter, dtype: int64
```

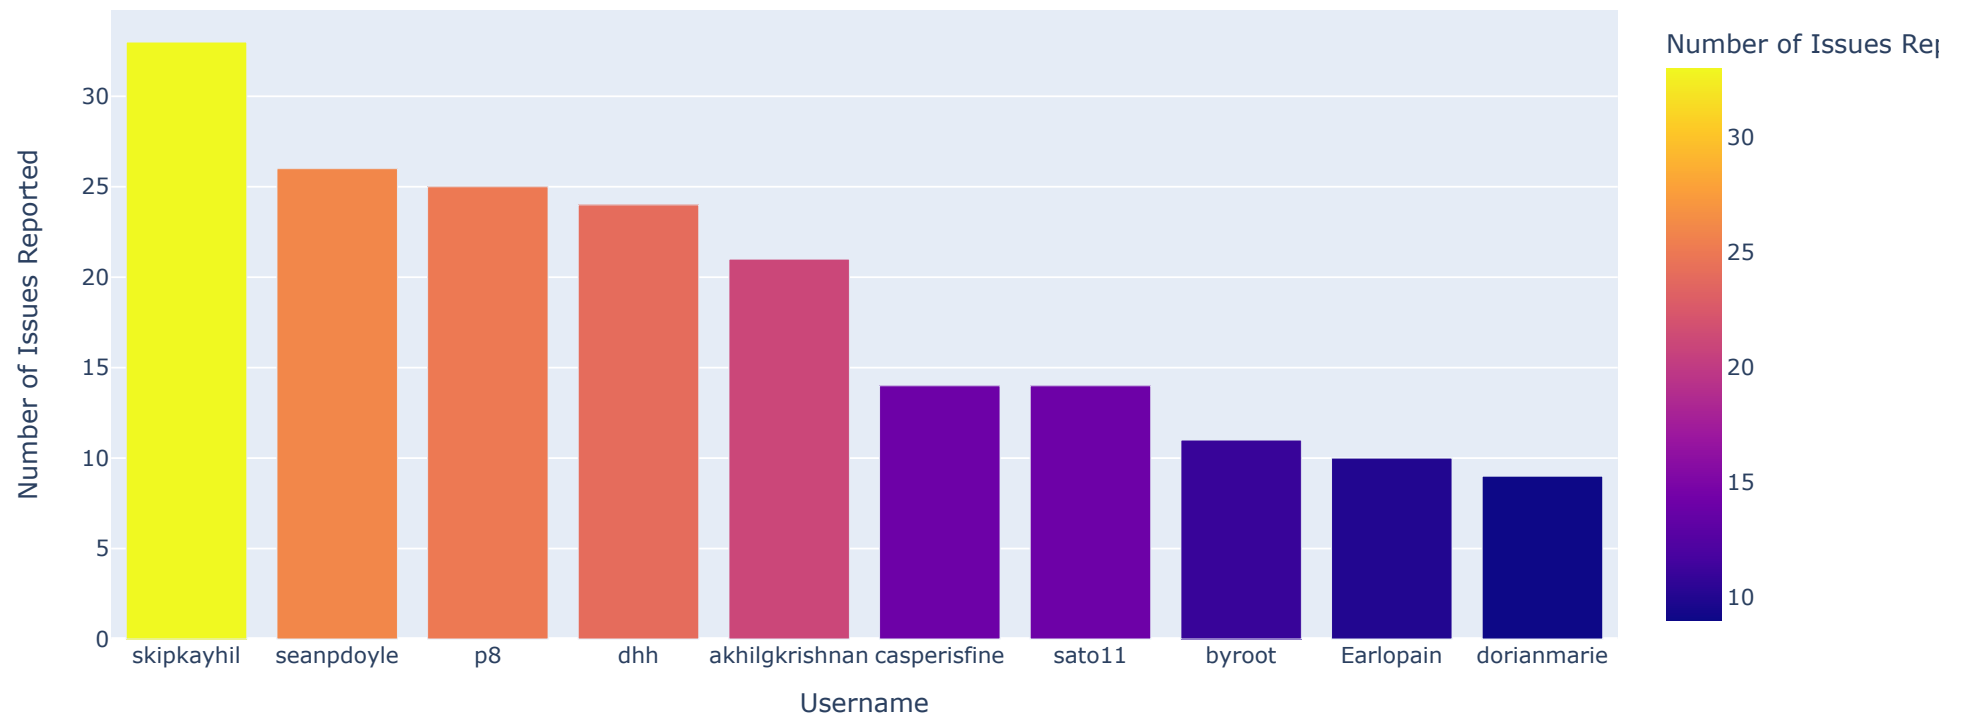
```
top_reporters_data = top_reporters.reset_index()
```

```
top_reporters_data.columns = ['Reporter', 'Frequency']
```

```
fig = px.bar(top_reporters_data, x='Reporter', y='Frequency',
             title='Top Issue Reporters',
             labels={'Reporter': 'Username', 'Frequency': 'Number of Issues Reported'},
             color='Frequency')
```

```
fig.show()
```

Top Issue Reporters



4. What is the most popular category (label)?

```
def extract_label_names(labels):
    return [label['name'] for label in labels if 'name' in label]

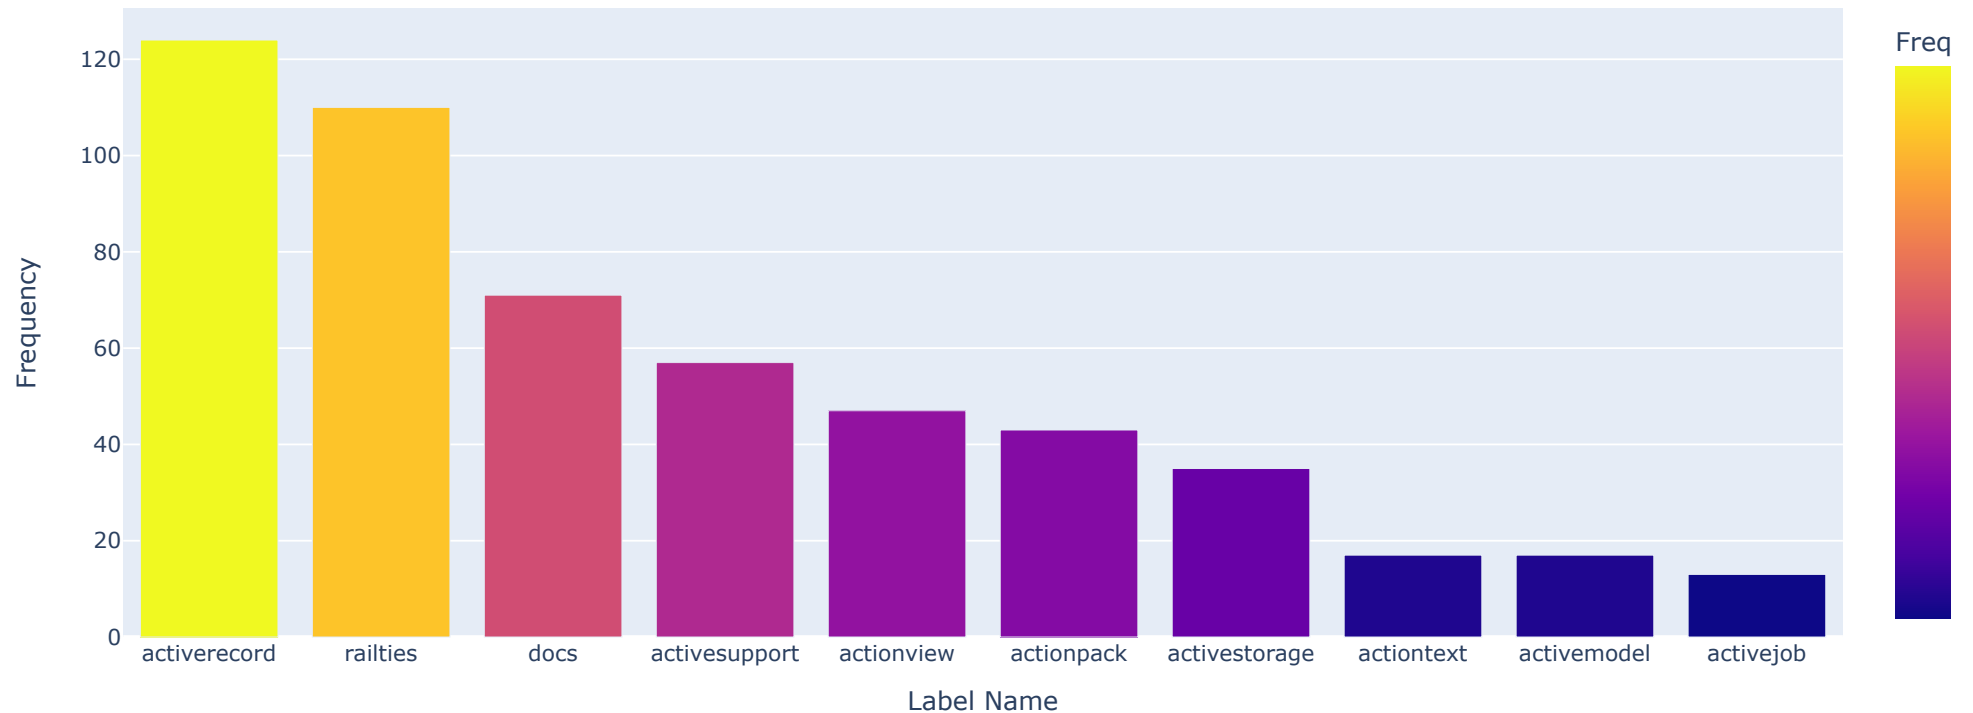
df['label_names'] = df['labels'].apply(extract_label_names)
all_label_names = sum(df['label_names'].tolist(), [])
label_counts = Counter(all_label_names)
labels_frequency_df = pd.DataFrame(label_counts.items(), columns=['Label Name', 'Frequency']).sort_values(by='Frequency', ascending=False).reset_

fig = px.bar(labels_frequency_df.head(10),
             x='Label Name',
             y='Frequency',
             title='Top 10 Most Popular Labels',
             labels={'Label Name': 'Label Name', 'Frequency': 'Frequency'},
             color='Frequency',
             )

fig.update_layout(xaxis_title="Label Name",
                  yaxis_title="Frequency",
                  xaxis={'categoryorder': 'total descending'})

fig.show()
```

Top 10 Most Popular Labels



▼ Most Discussed Issues

```
most_discussed_issues = df.sort_values(by='comments', ascending=False)
print(most_discussed_issues[['title', 'comments']].head(10))
```


	title	comments
496	Set a new default for the Puma thread count	83
490	Add (a very basic!!) Rubocop by default	30
180	Request: rename Rails console as irr or someth...	23
384	Generate devcontainer files by default	20

467	Warning on serialize from ActiveStorage	13
492	Extract Action Notifier framework for push not...	13
439	Default to creating git pre-commit	13
460	Add rate limiting to Action Controller via the...	13
63	Introduce `ActiveSupport::TestCase.around`	11
471	Add Thruster to Docker setup to get HTTP/2, X-...	11

```
issue_counts_by_month = df.resample('M', on='created_at').size()
```

```
# Plot
fig = px.line(issue_counts_by_month, title='Number of Issues Each Month')
fig.update_xaxes(title_text='Month')
fig.update_yaxes(title_text='Number of Issues')
fig.show()
```


Number of Issues Each Month



Classification Task

```

import pandas as pd
import torch
from torch.utils.data import Dataset, DataLoader
from sklearn.model_selection import train_test_split
from transformers import AutoTokenizer, AutoModelForSequenceClassification, AdamW, get_linear_schedule_with_warmup, Trainer, TrainingArgument
from sklearn.metrics import accuracy_score, precision_recall_fscore_support
import numpy as np

def get_first_label(label_list):
    if label_list:
        return label_list[0]['name'] # This gets the name of the first label
    else:
        return 'No Label'

df['single_label'] = df['labels'].apply(get_first_label)

# Assuming df is your DataFrame
df['body'] = df['body'].astype(str) # Ensure text column is string
df['single_label'] = pd.factorize(df['single_label'])[0]

class GitHubIssueDataset(Dataset):
    def __init__(self, encodings, labels):
        self.encodings = encodings
        self.labels = labels

    def __getitem__(self, idx):
        item = {key: torch.tensor(val[idx]) for key, val in self.encodings.items()}
        item['labels'] = torch.tensor(self.labels[idx])
        return item

    def __len__(self):
        return len(self.labels)

# Load tokenizer
tokenizer = AutoTokenizer.from_pretrained('bert-base-uncased')

# Prepare datasets
X_train, X_val, y_train, y_val = train_test_split(df['body'], df['single_label'], test_size=0.2)
train_encodings = tokenizer(X_train.tolist(), truncation=True, padding=True)
val_encodings = tokenizer(X_val.tolist(), truncation=True, padding=True)
train_dataset = GitHubIssueDataset(train_encodings, y_train.tolist())

```



```

val_dataset = GitHubIssueDataset(val_encodings, y_val.tolist())

# Load model
model = AutoModelForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=len(df['single_label'].unique()))

# Training arguments
training_args = TrainingArguments(
    output_dir='./results',
    num_train_epochs=3,
    per_device_train_batch_size=8,
    warmup_steps=500,
    weight_decay=0.01,
    logging_dir='./logs',
    evaluation_strategy="epoch",
    logging_steps=10,
    save_strategy="epoch",
    load_best_model_at_end=True,
)

# Custom compute_metrics function
def compute_metrics(pred):
    labels = pred.label_ids
    preds = pred.predictions.argmax(-1) # Convert model logits to class predictions
    # Use 'macro', 'micro', or 'weighted' averaging based on your specific needs
    precision, recall, f1, _ = precision_recall_fscore_support(labels, preds, average='macro')
    acc = accuracy_score(labels, preds)
    return {
        'accuracy': acc,
        'f1': f1,
        'precision': precision,
        'recall': recall
    }

# Initialize Trainer
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=val_dataset,
    compute_metrics=compute_metrics,

```

)

```
# Train the model  
trainer.train()
```